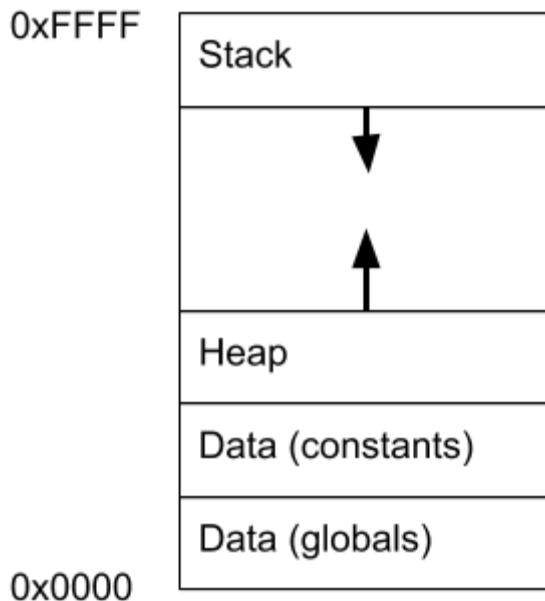


Memory Management




- Local variable allocated on the stack
 - Lifetime: when it's called → when the function ends
- Global variables put in the data segment
 - Automatically allocated by the OS
 - Lifetime: entire program
 - There is also a read-only data segment for constants
- The heap
 - Programmer-managed area of memory
 - Has nothing to do with the heap data structure
 - You can create and destroy pieces of memory on demand
 - Heap lifetimes can cross boundaries
- Stack and the heap both change size
 - One end of the stack is fixed, one end changes
 - Stack grows down (high addresses to lower addresses)
 - Heap grows up (lower addresses to higher addresses)
 - If they meet, the program is out of memory
 - Why is it arranged this way?
 - Arbitrary convention
 - Why not grow outward from the center?
 - You would be basically throwing away half of the memory
 - The heap would only be able to occupy half the memory
 - Why in the same area of memory at all?
 - If you have a 16/32 bit computer, you don't have any options

Heap functions

- `<stdlib.h>`
- `malloc`
 - To allocate memory

- `int* arr = malloc(sizeof(int) * 20;`
 - Gives space for 20 ints on the heap assigned into a pointer named arr




- `malloc` takes the number of bytes to allocate, makes a block of bytes at least that big, then returns a pointer to the block
- `free`
 - When you're done with that block of memory, use `free`
 - `free(arr);`
 - If you `malloc` a second thing, it might use that freed data
 - But there's still a pointer in the stack
- Void pointers
 - `void* malloc(size_t size);`
 - An `int*` points to an `int`, a `FILE*` points to a `FILE` object
 - A `void*` is a universal pointer
 - It can point to anything
 - Any address can go into `void*`
 - A memory address with no meaning
 - Can be assigned any other type of pointer
 - `int x; float y;`
 - `void* vp = &x; // ok`
 - `vp = &p; // ok`
 - Can't access data from a `void*` without casting

Allocating a struct on the heap

- Similar to `new` in Java
- `food* f = malloc(sizeof(food));`
- Like anything on the stack, everything you `malloc` contains garbage
 - Papers on the stack of papers might have writing on them
 - Need to allocate an initialized struct to avoid undefined behavior
- To fix this, you can use `memset` from `<string.h>`
 - `memset(f, 0, sizeof(food)); // will fill the memory at f of size of food number of bytes with 0s`
- Could also use `calloc`
 - `food* f = calloc(sizeof(food), 1);`

- `calloc` is a `malloc` followed by a `memset`
- “Allocated `1*sizeof(food)` bytes and fill them with 1s”
- Zeroing out all the data fields so they’re not filled with garbage
- It’s faster to `malloc` and initialize yourself
- Caveats:
 - Do NOT do:
 - `while(1) malloc(1048576);`
 - Program will run out of memory without error; it loops forever because if you run out of memory, `malloc` returns `NULL`
 - `int* arr = malloc(sizeof(int)*20);`
 - `free(arr);`
 - `arr[0] = 100;`
 - As soon as you free pointers, all pointers to it become invalid and have undefined behavior
 - It’s your responsibility as the user to not use invalid pointers
 - Can only `free` a piece of memory ONCE



Heap rules:

1. When you `malloc` you should check if it returns `NULL`
2. You must `free` everything you allocate
3. You must only `free` it exactly once
4. You must not access memory that has been freed
 - a. Who is responsible for deallocating? How do you know how many things are using a piece of memory? What if nobody points to it?

VLAs in C99

- Modern c code can allocate variable-sized pieces of memory on the stack instead of the heap
 - `int len = strlen(str);`
 - `char newArray[len + 1];` //not a constant size
- Called a variable-length array (VLA)
- It’s still on the stack
- Increases the size of the activation record
- Gives automatic lifetime
 - Automatically freed

- Not used for everything because stack space is limited
 - On Thoth, it's only 10 MB
- Also doesn't solve the problem of ownership
 - Sometimes you need data to stick around longer than the length of the function
 - Have to hand the pointer off
 - Can't return pointers as local variables
- Also stack is linear, cannot implement a tree on the stack

Garbage collection:

- If nothing is pointing at something and the something isn't pointing at anything, it gets garbage collected
- Roots: any part of memory that isn't the head (stack, globals, etc.)
 - Program is using it currently, cannot get rid of it
- Reachability: an object is reachable if there is a path from the roots to the object
- When you remove the stack variable, the only way to reach the object, it becomes garbage
- In C, there is no garbage collector and the heap can get filled up
 - This is a memory leak: losing the last reference to a piece of memory
 - No link to an object from the roots
 - We can never deallocate it, it leaked out of our program
 - The only way to deallocate it is to exit your program
 - All of your program's memory is deallocated when you exit
 - Never take garbage collection for granted